Sequence Encoding

Multi-Head Attention

國立臺灣大學 資訊工程學系

陳縕儂 助理教授

http://vivianchen.idv.tw

Convolutions

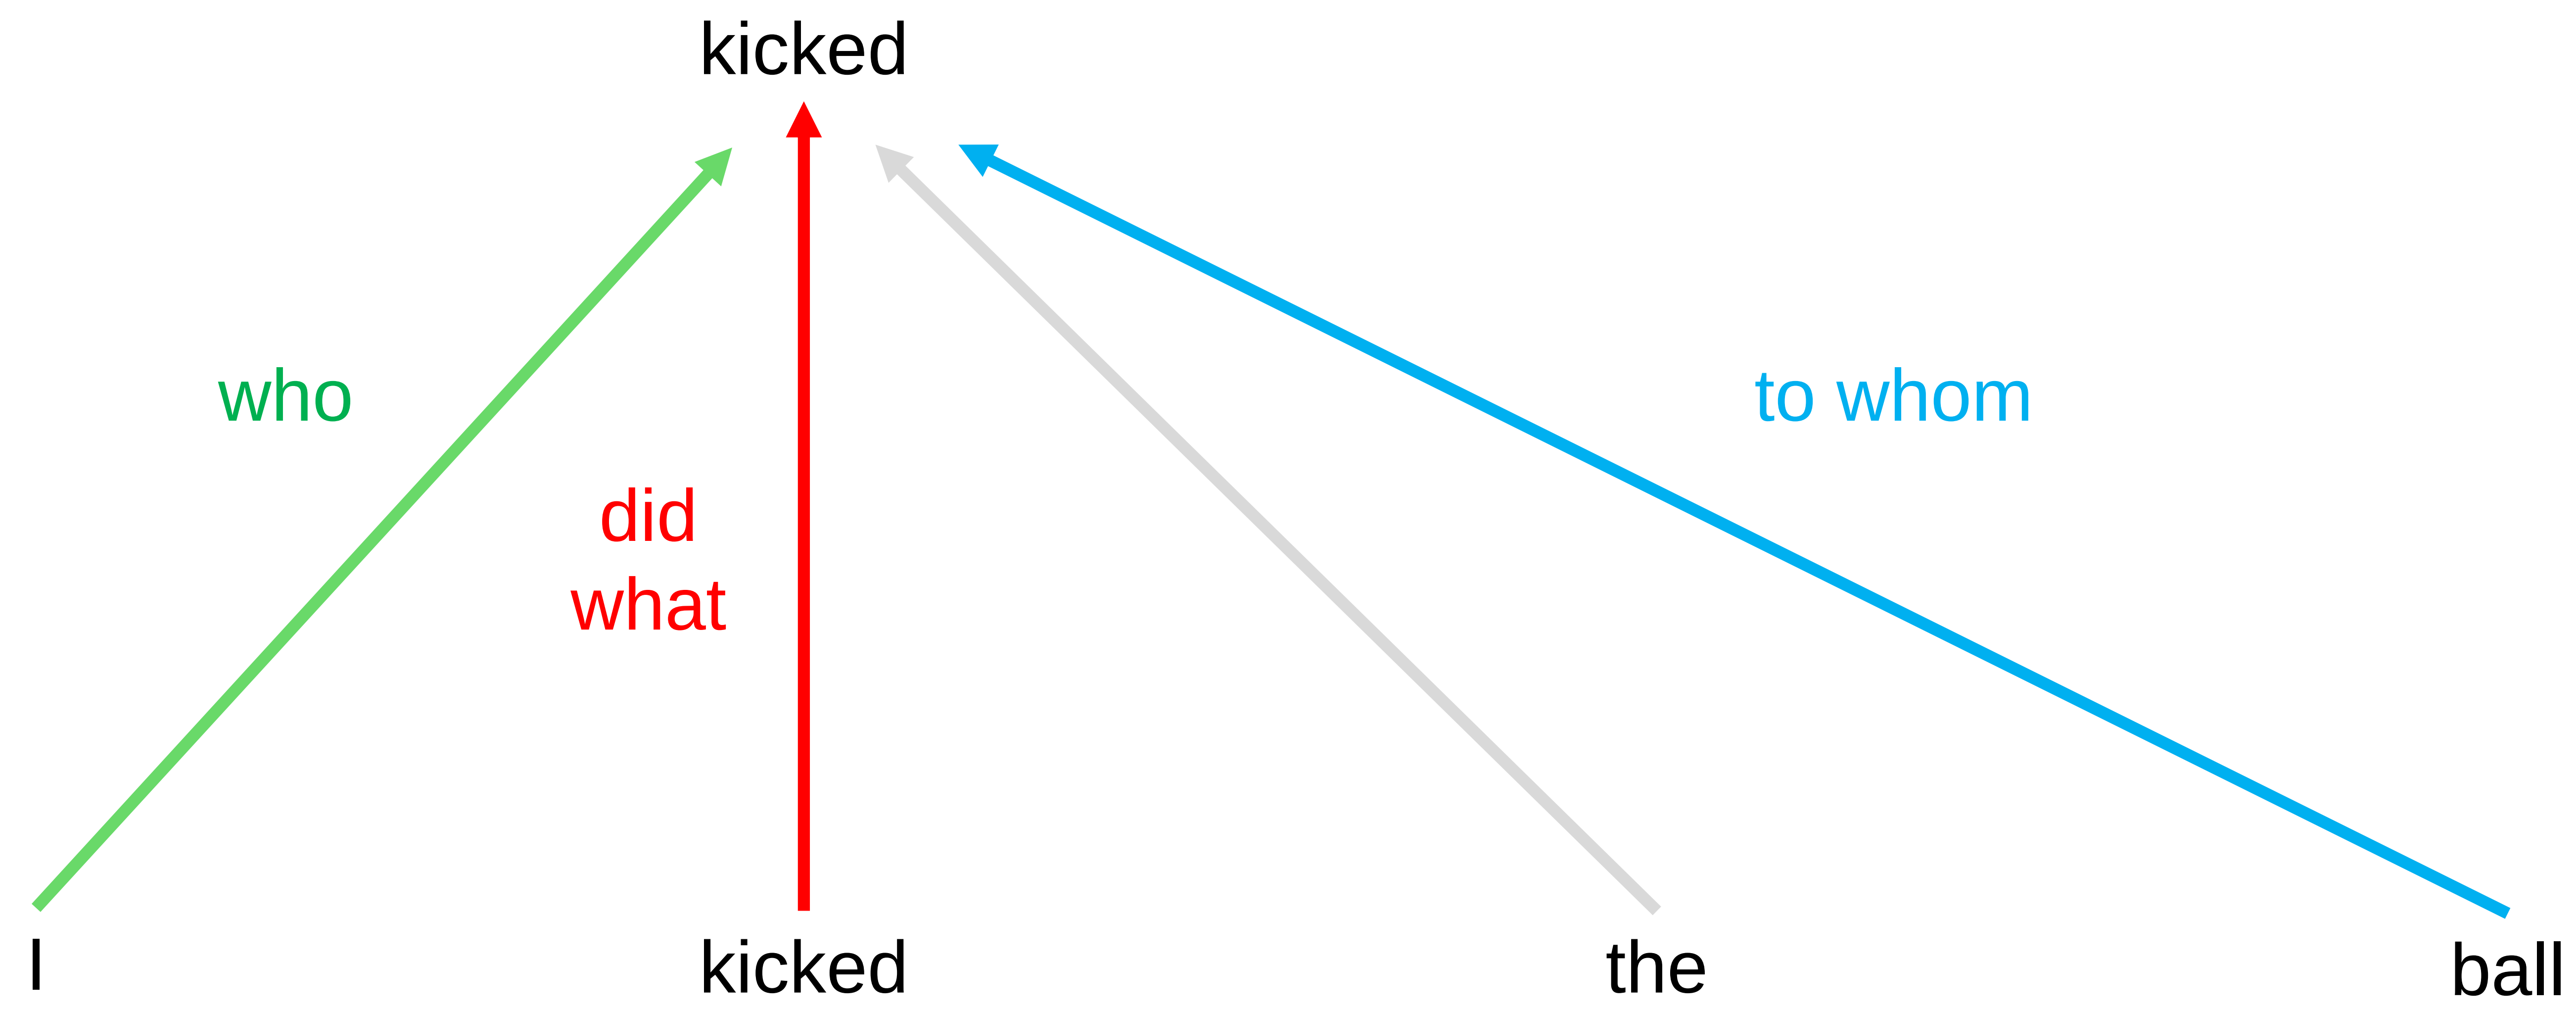

Self-Attention

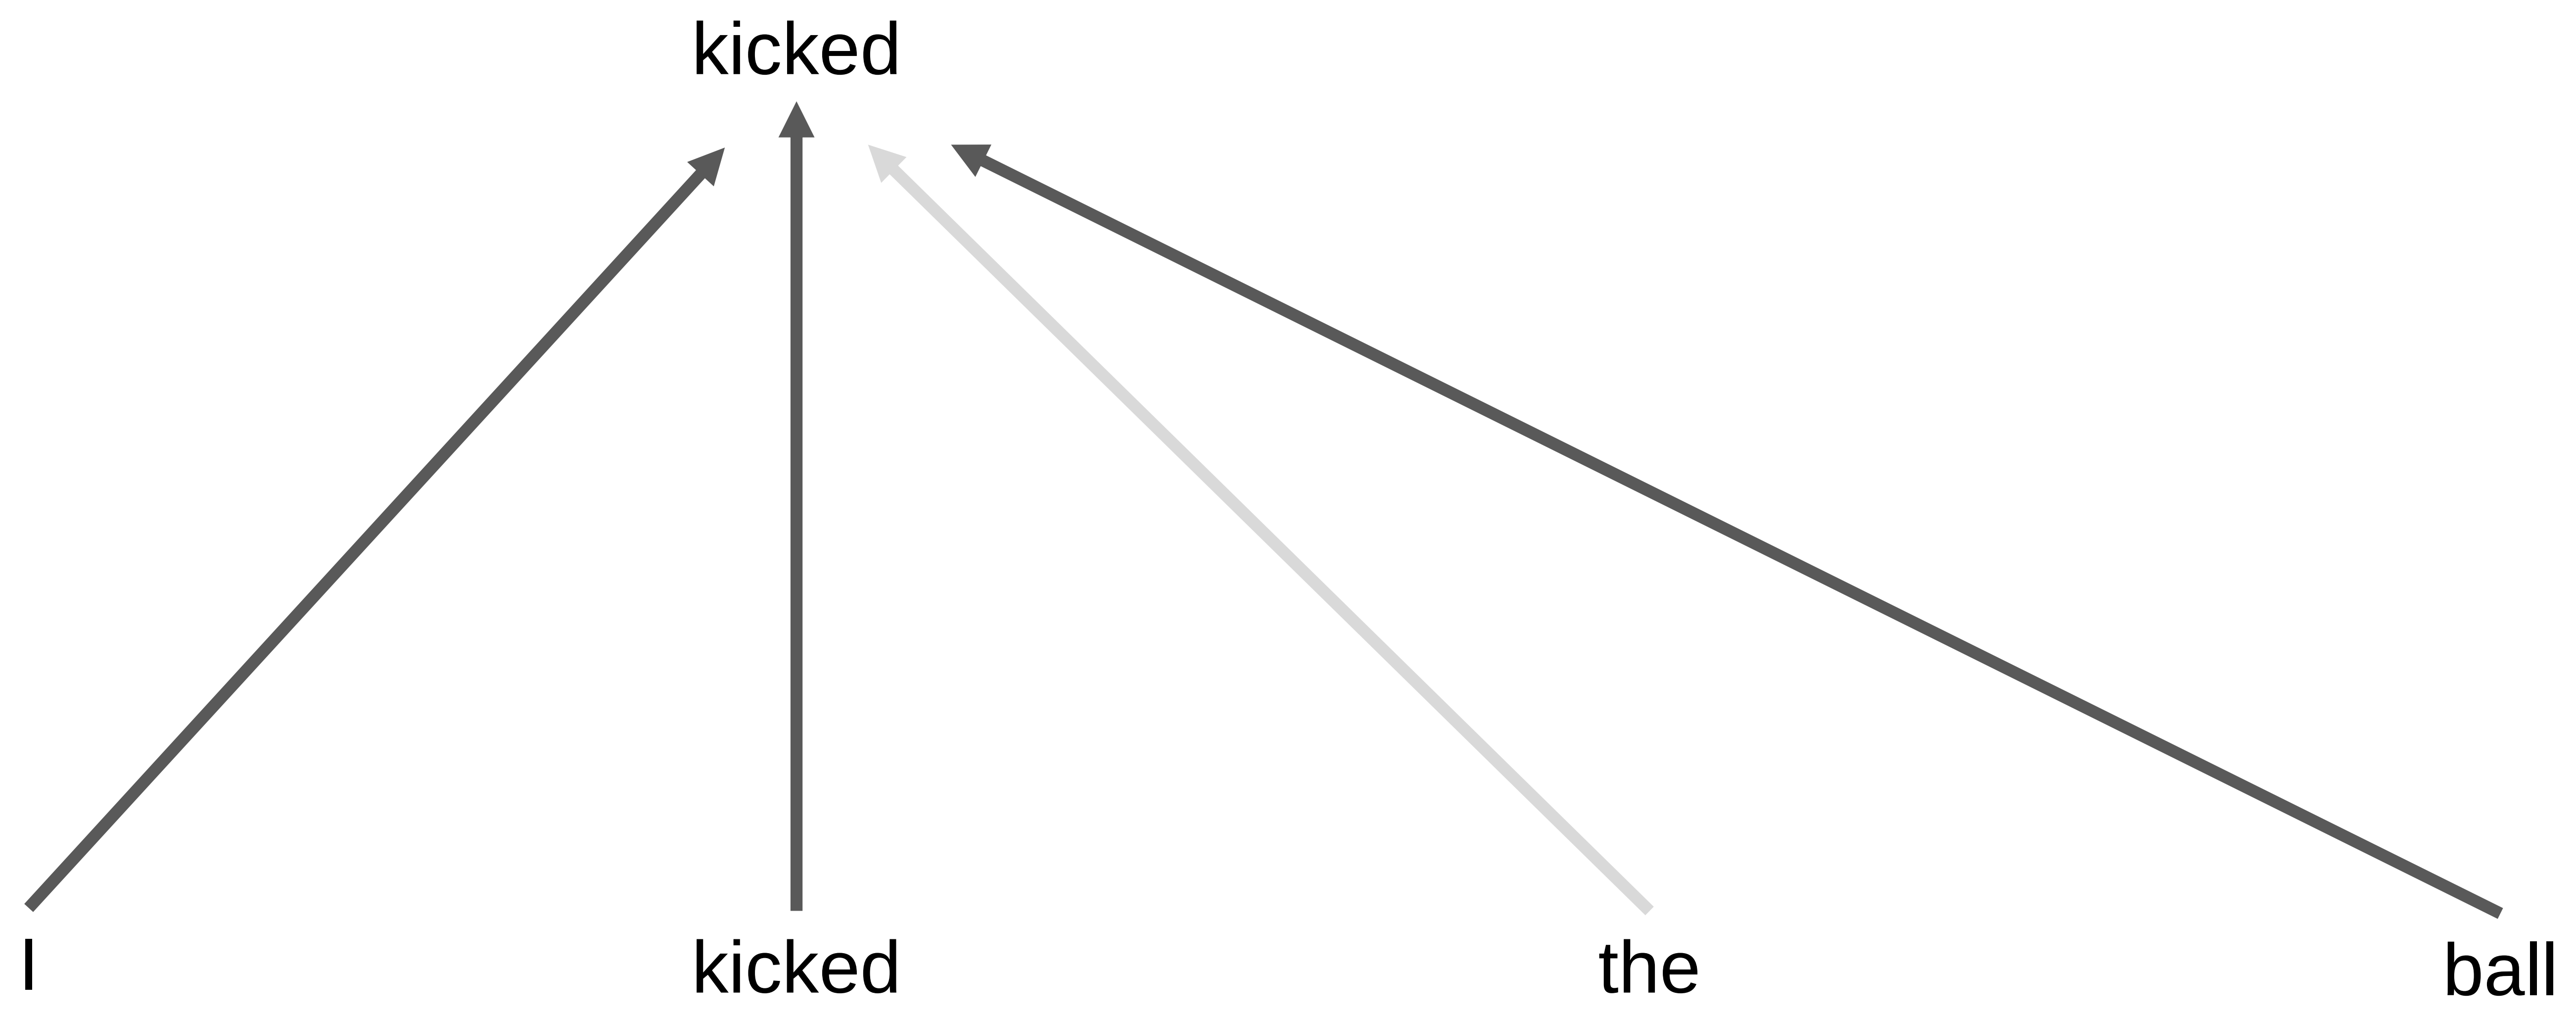
kicked

I

kicked

the

ball

Attention Head: who

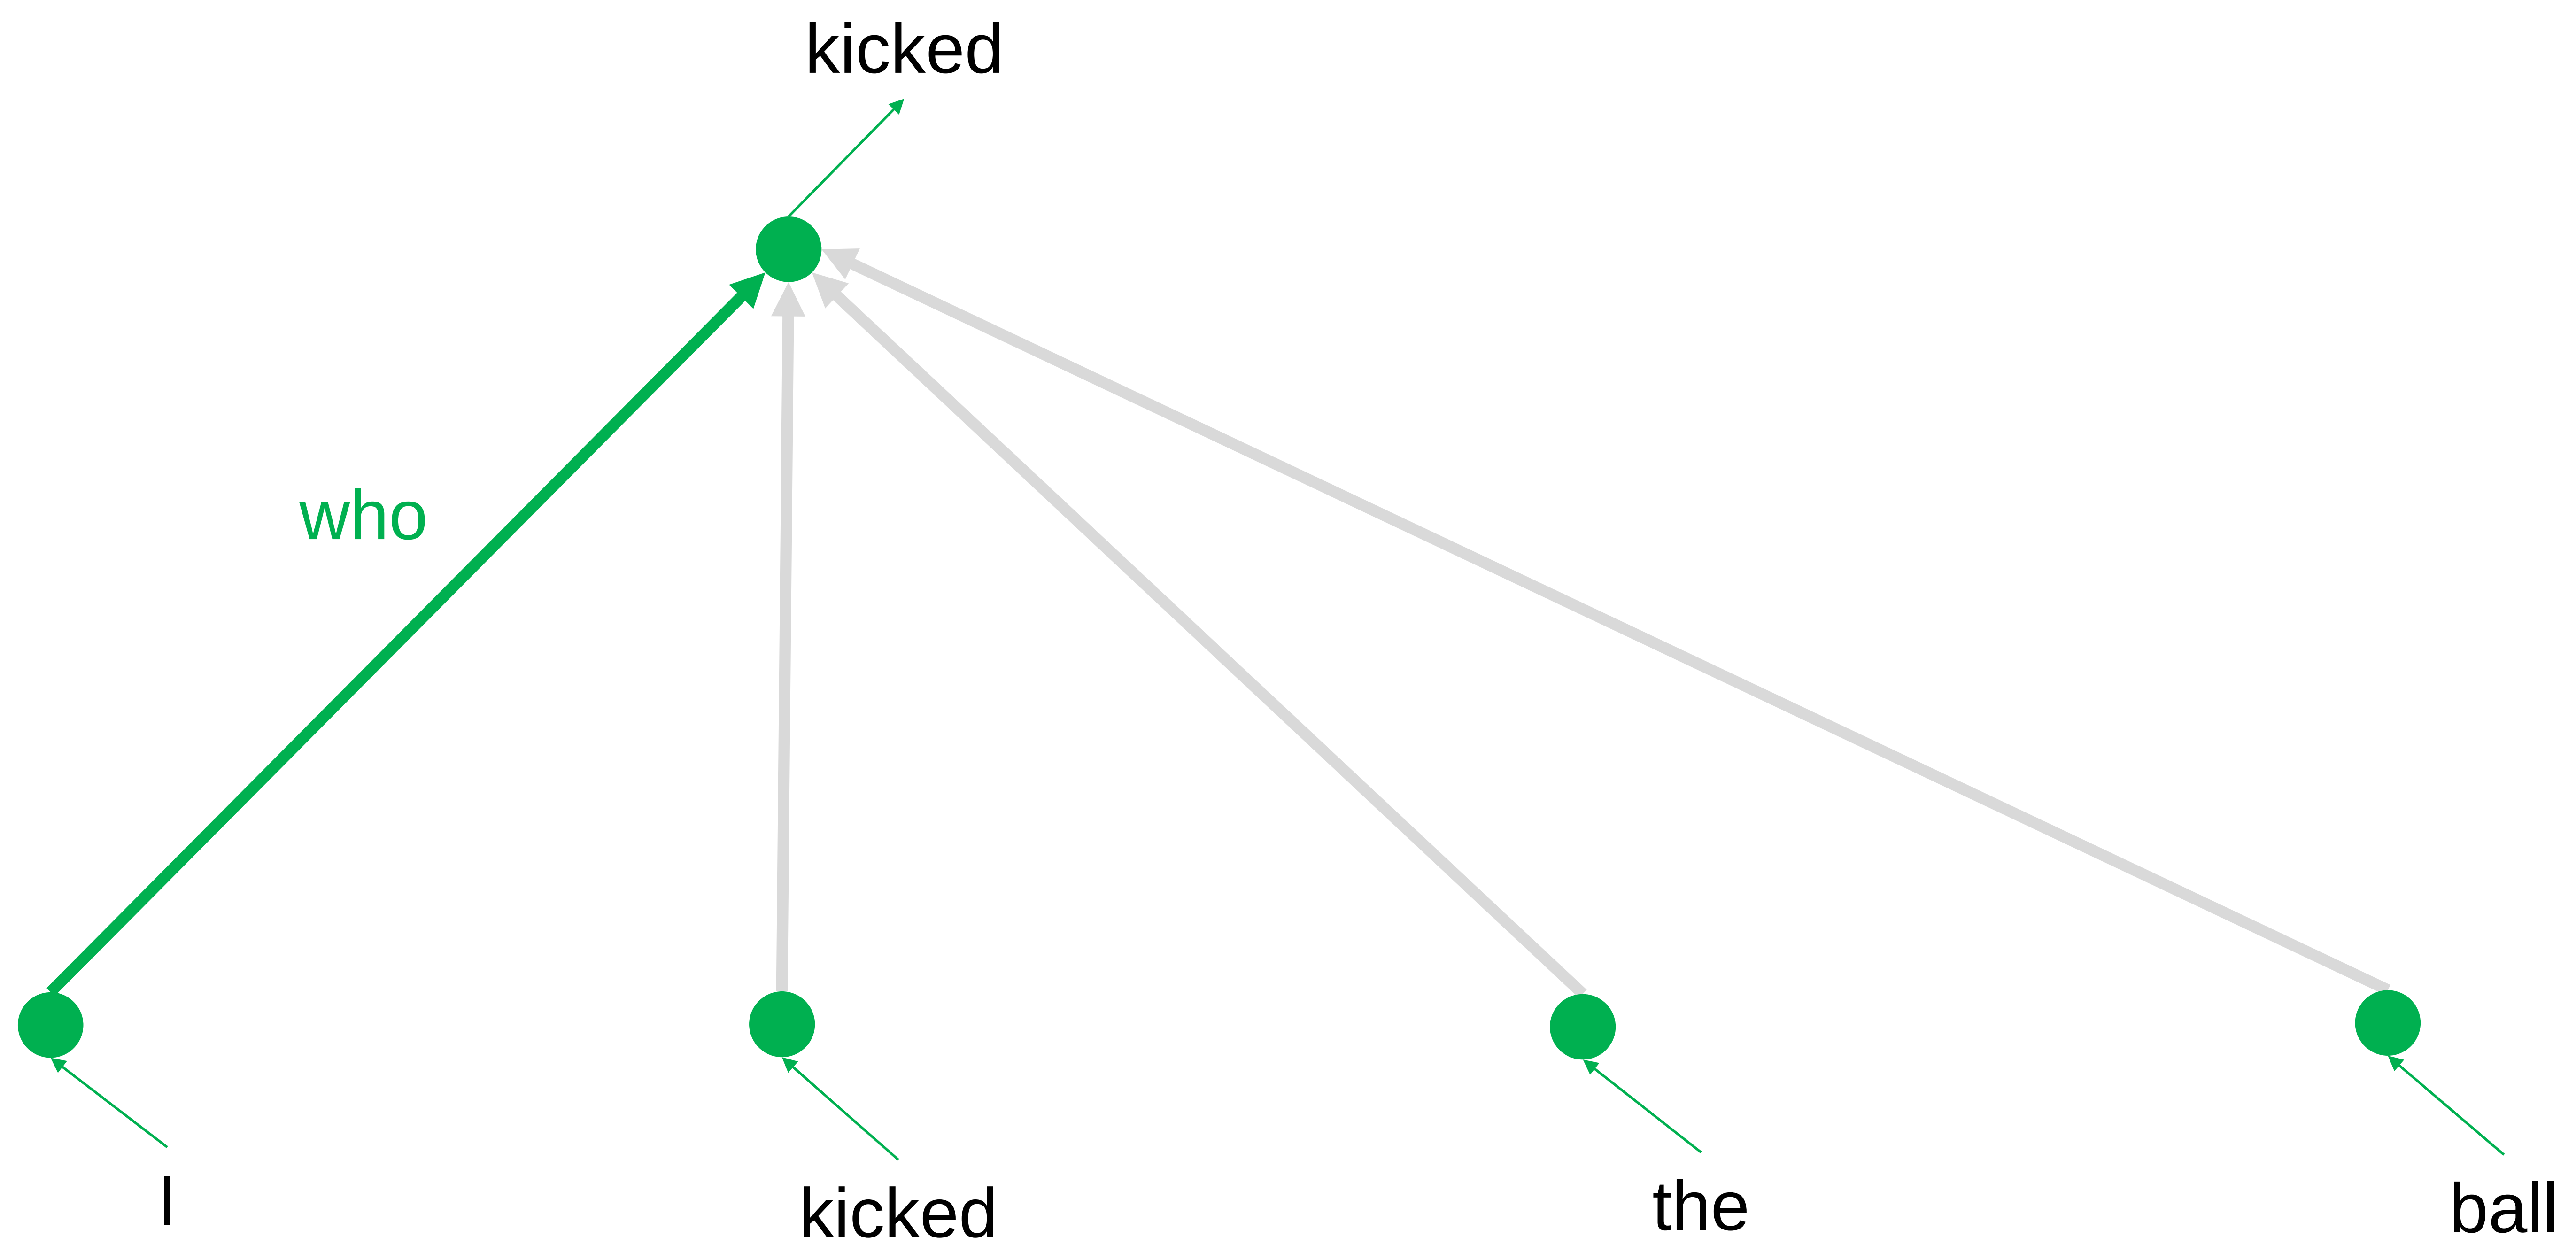

Attention Head: did what

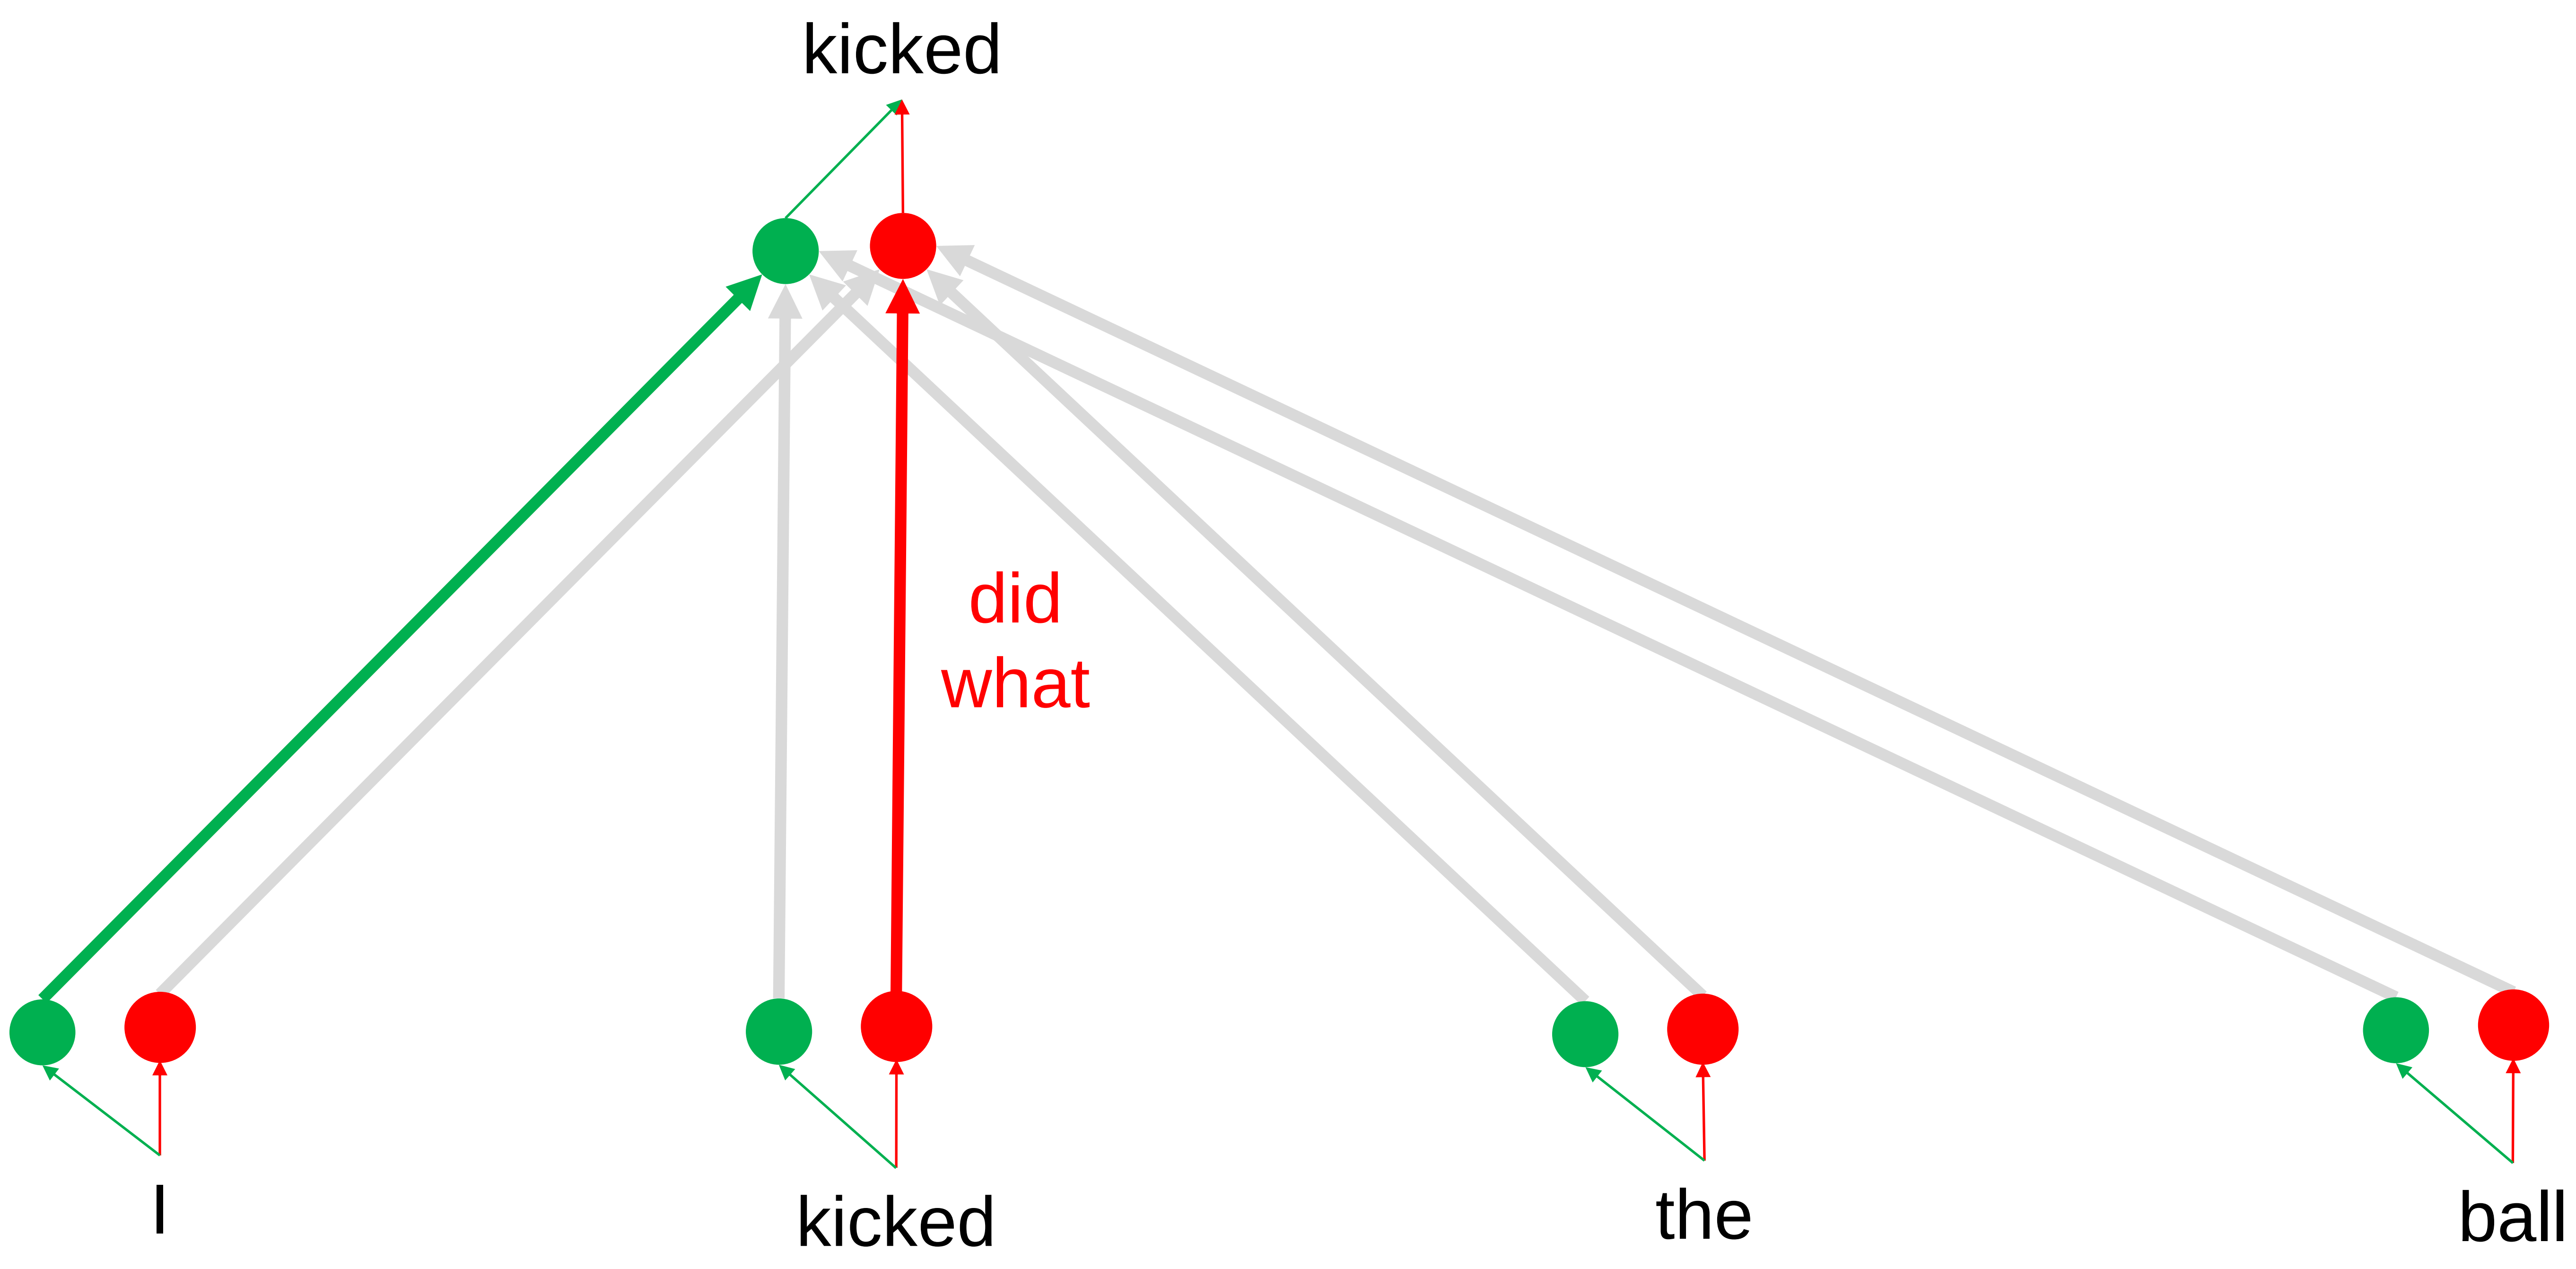

Attention Head: to whom

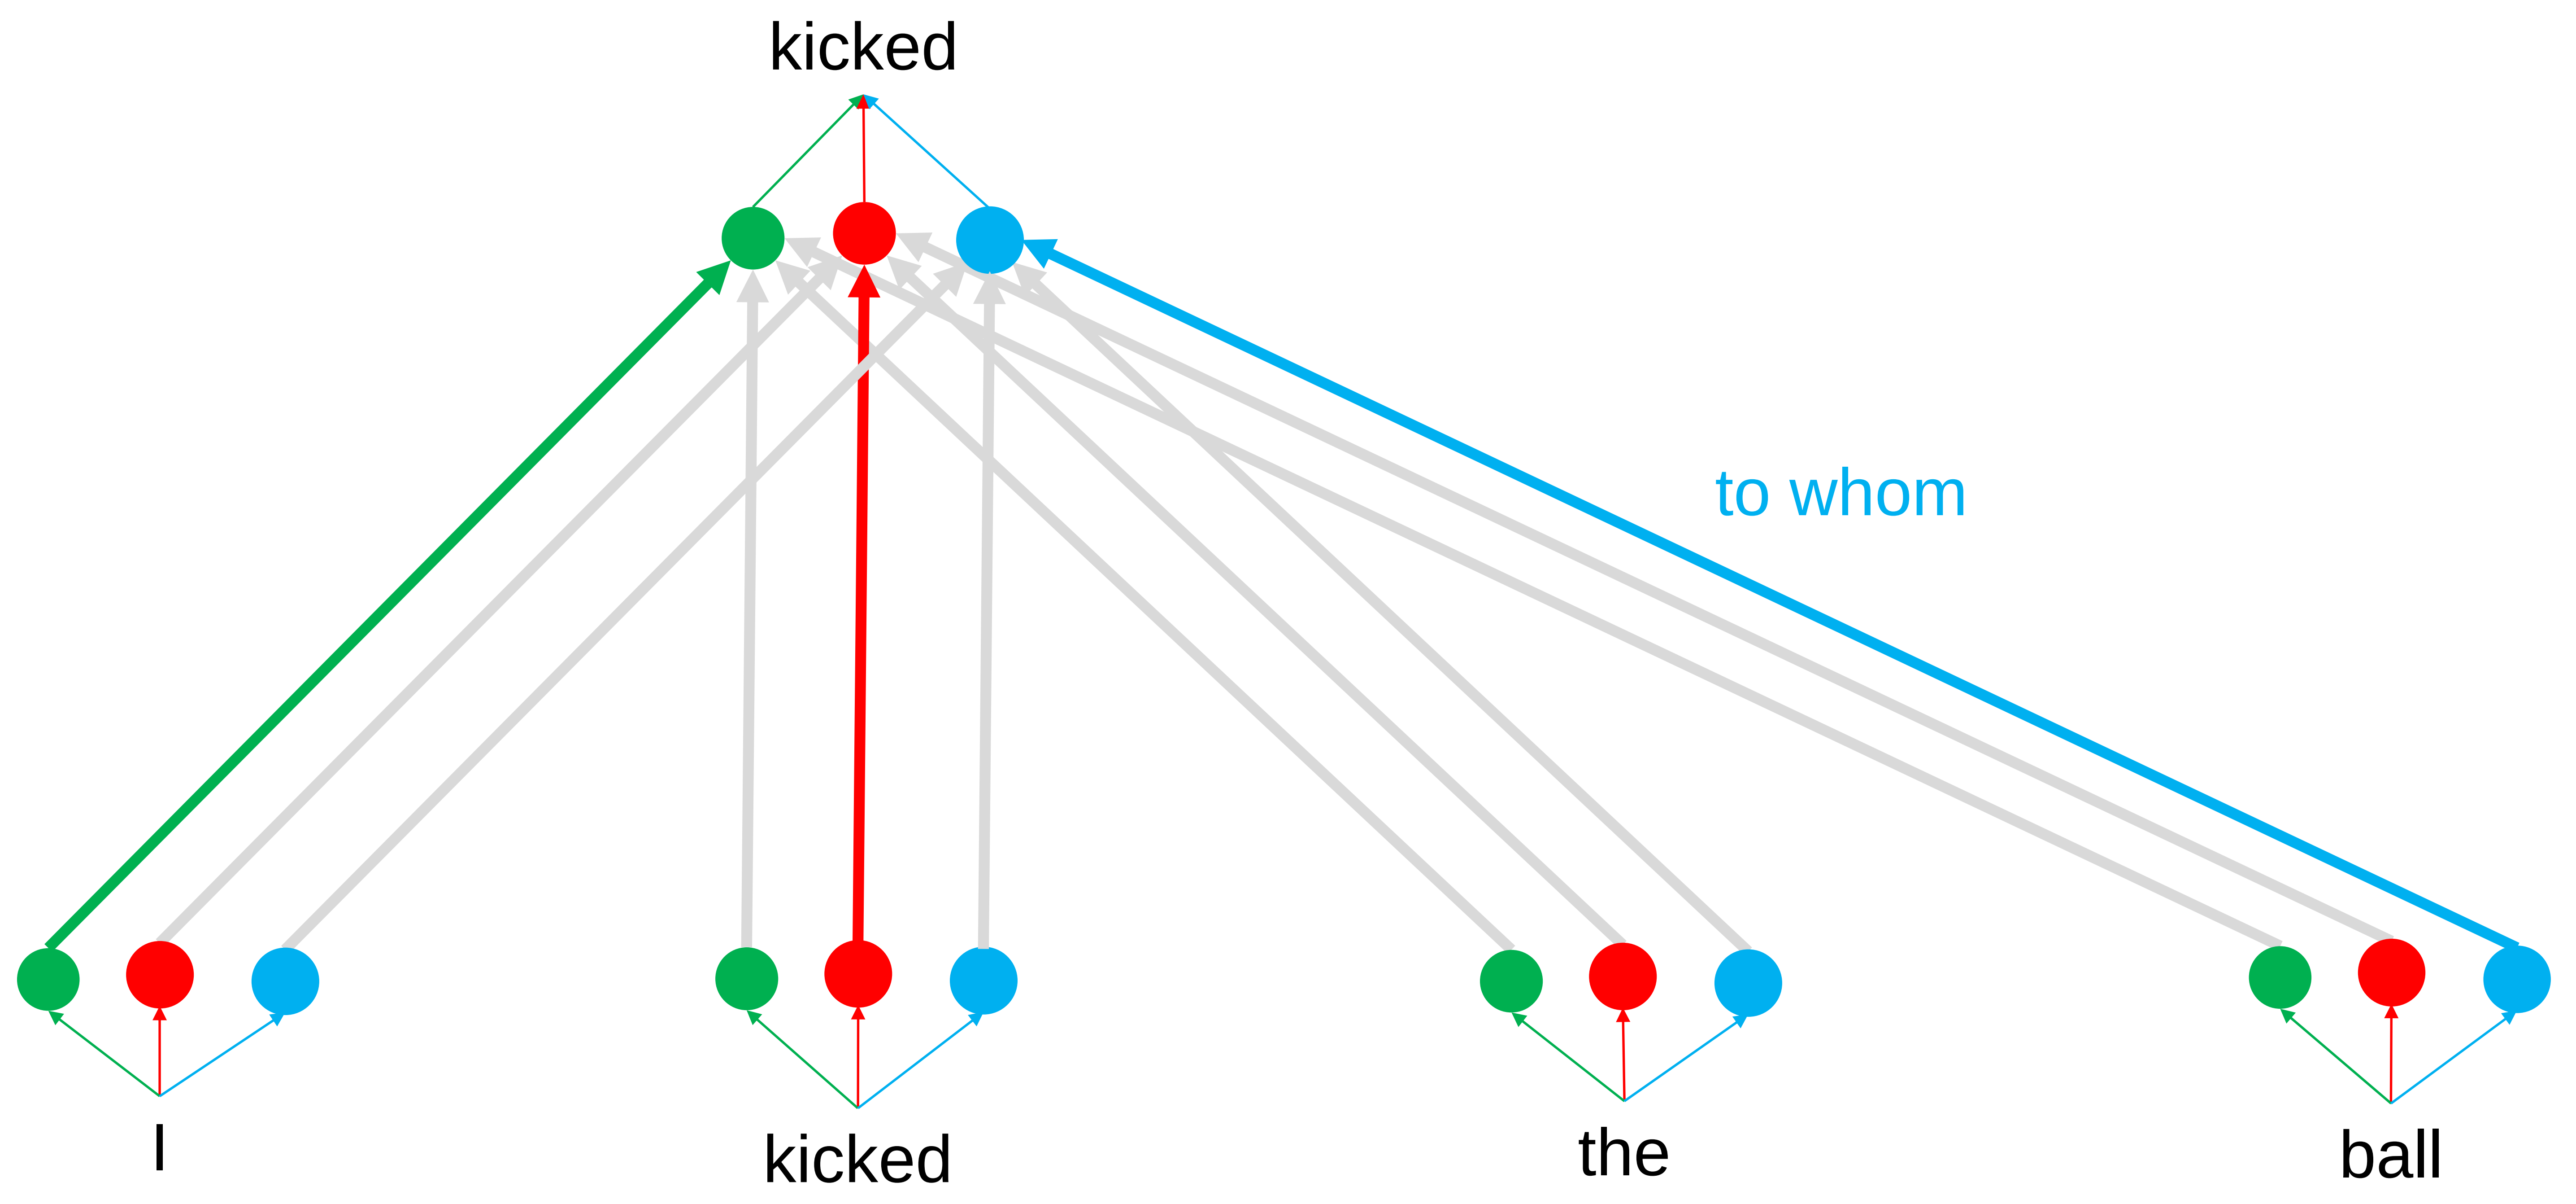

Multi-Head Attention

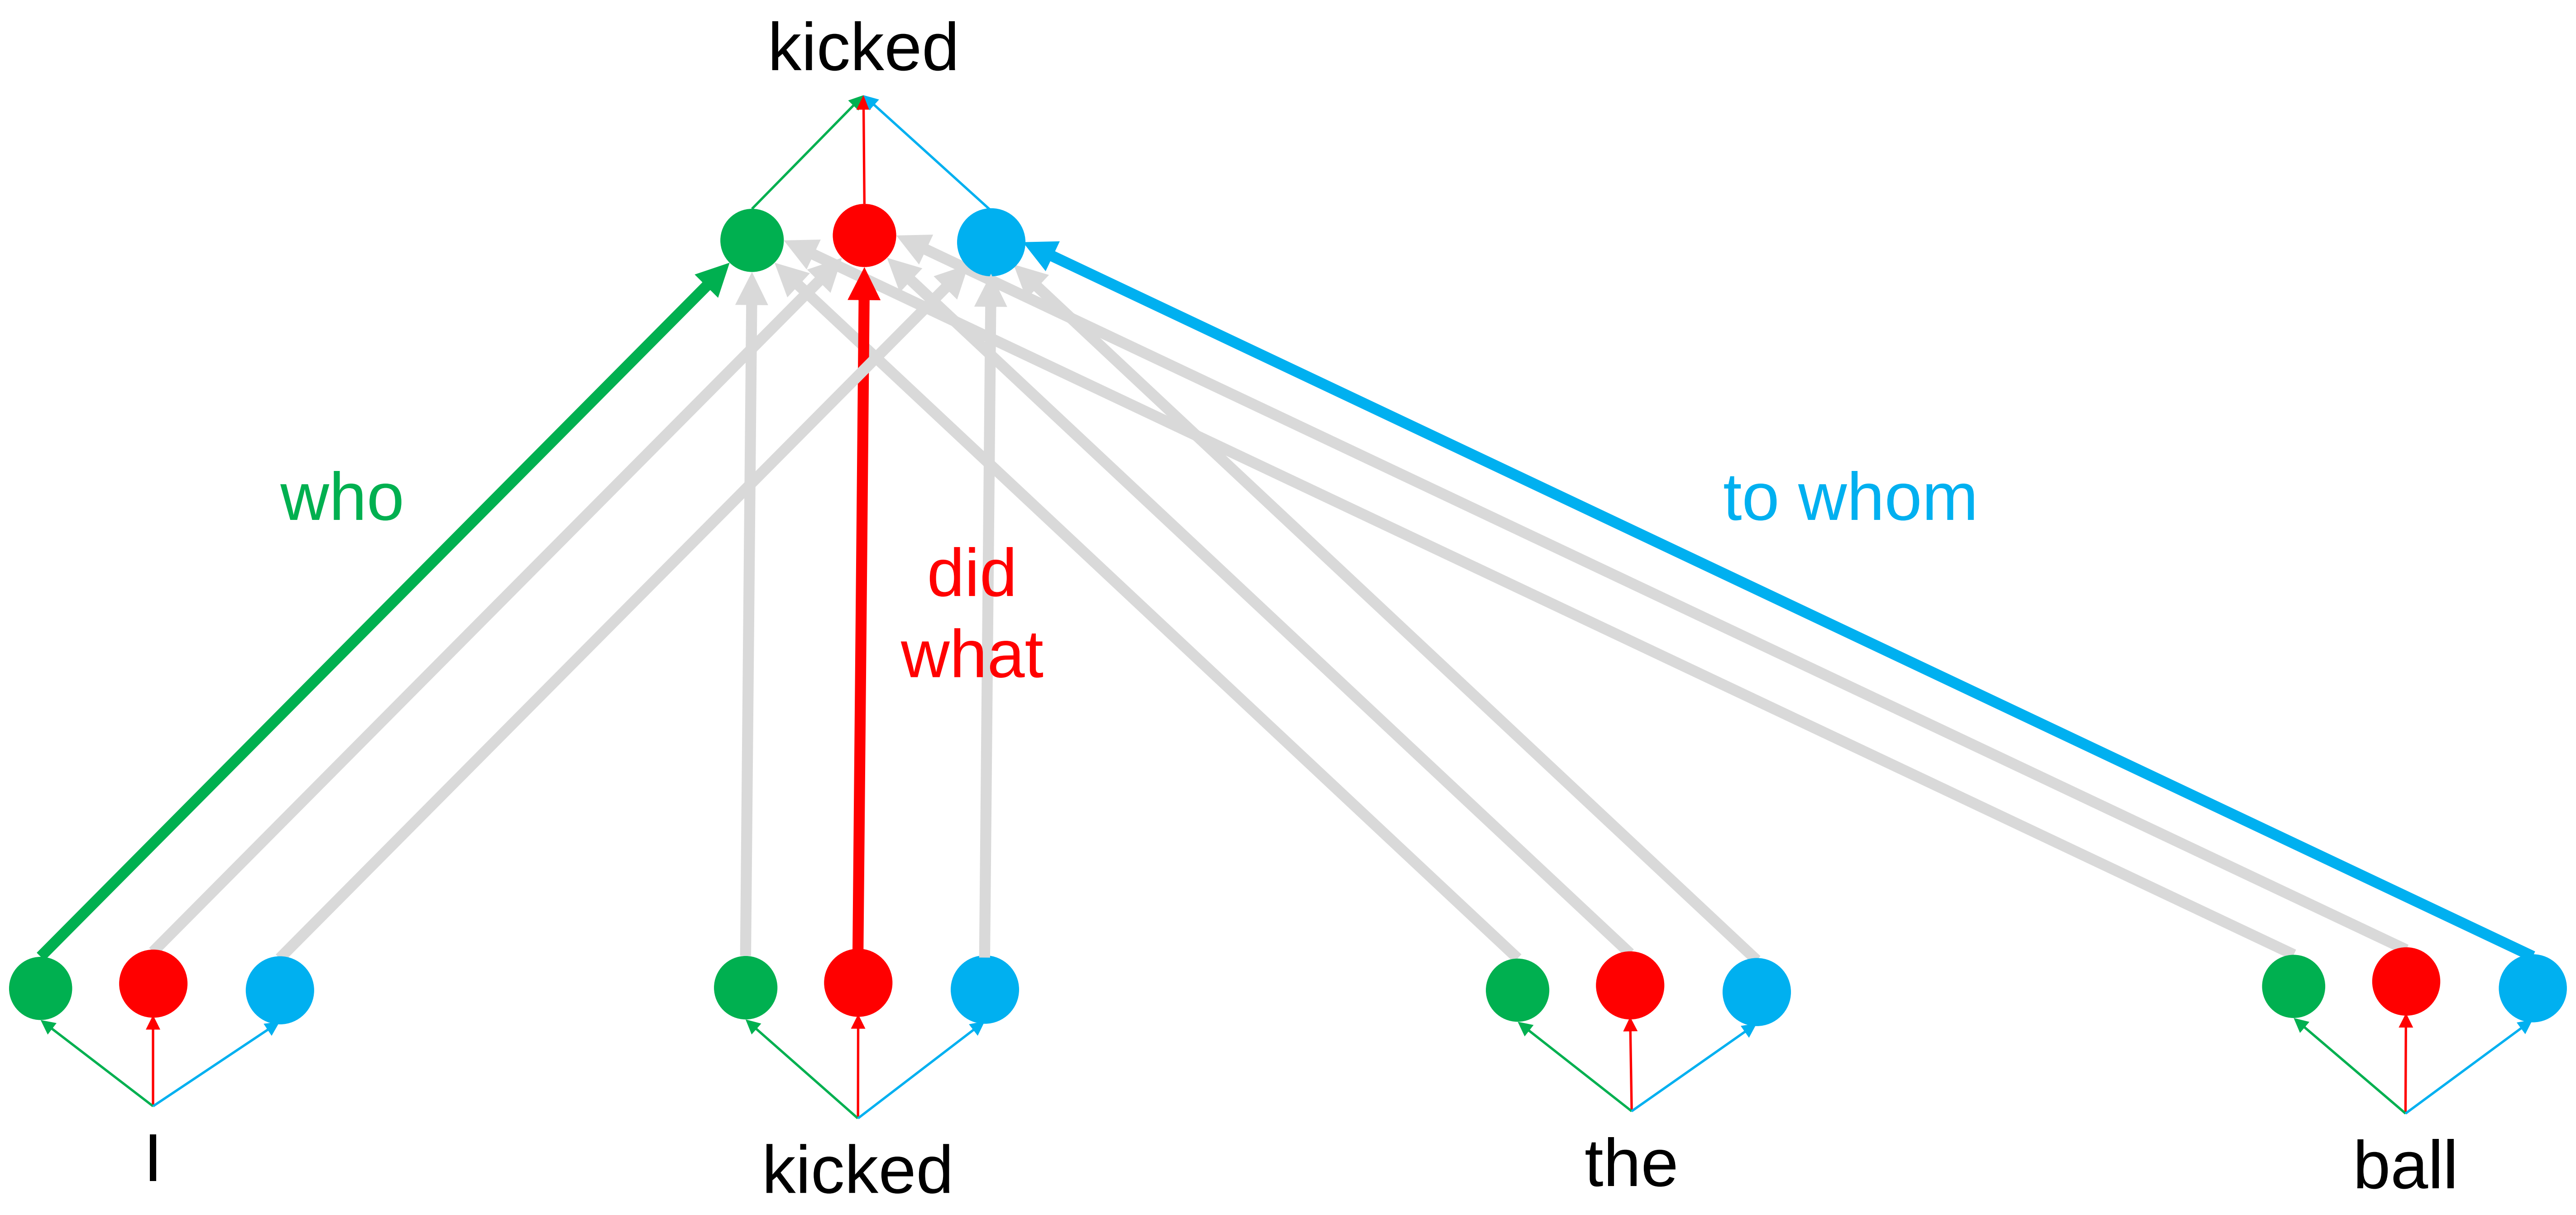

Comparison

- Convolution: different linear transformations by relative positions

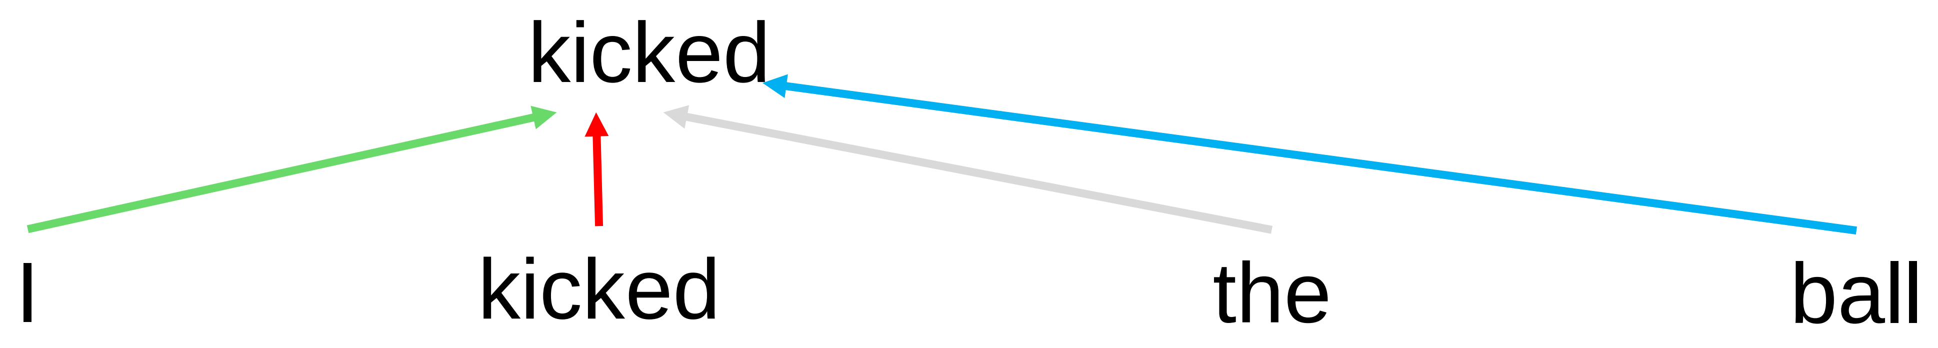

- Attention: a weighted average

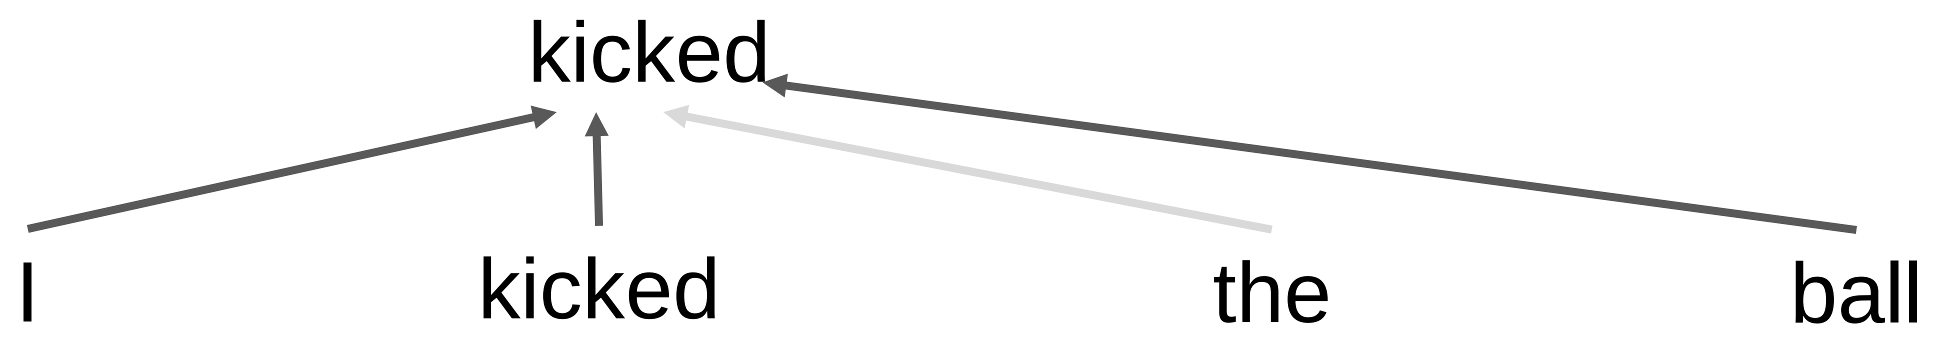

- Multi-Head Attention: parallel attention layers with different linear transformations on input/output

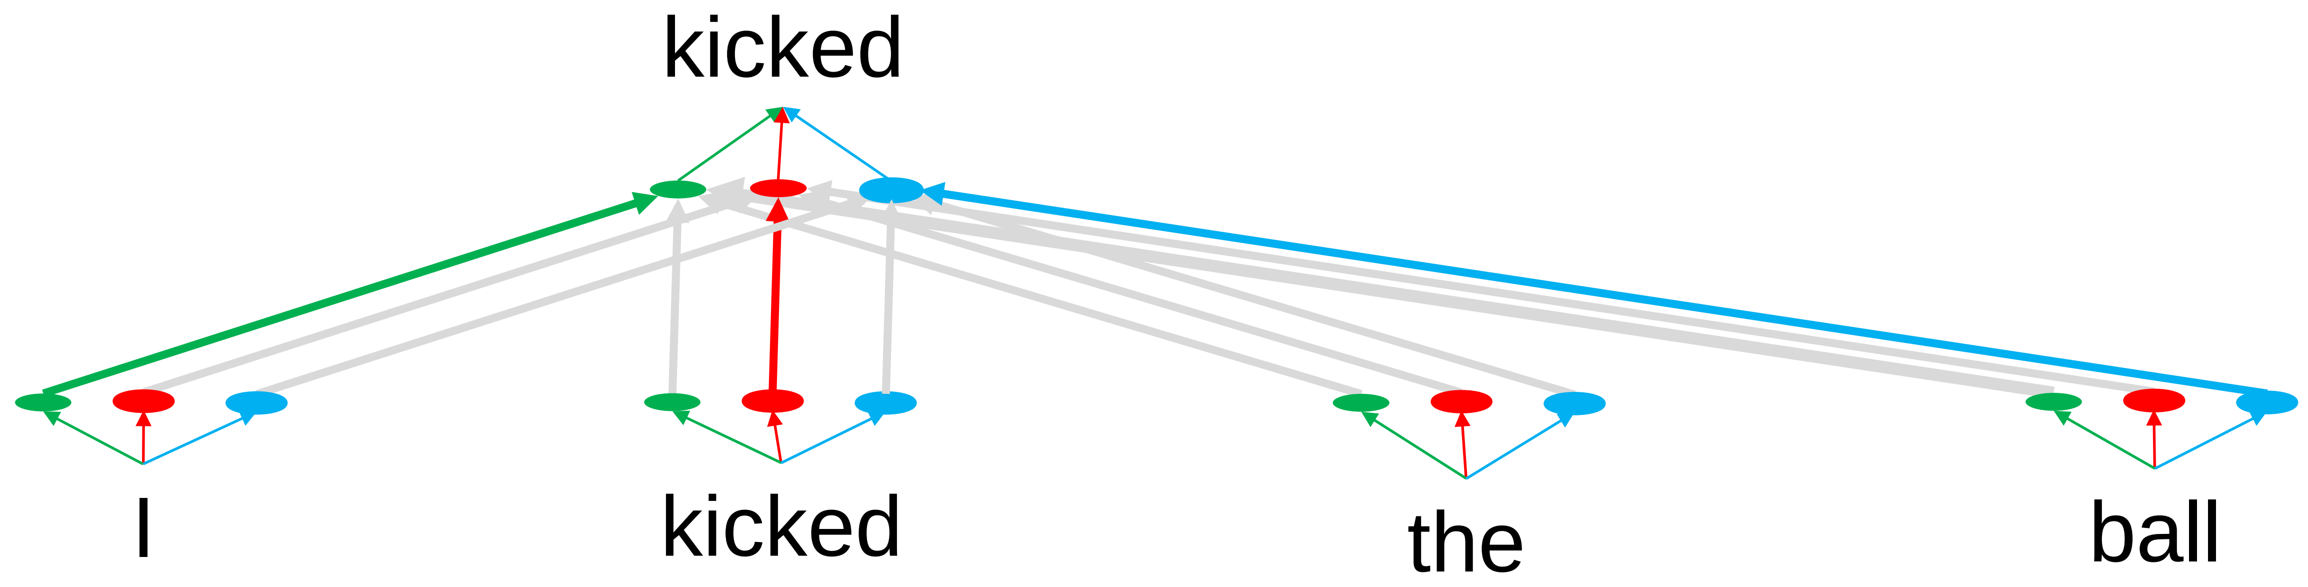
kicked
I
kicked
the
ball